

Assignment-A2 41403

- Title: Parallel computing using CUDA
- Problem statement: Vector and matrix operations
design parallel algorithm to
 - add two large vectors
 - multiply vector & matrix
 - multiply two $N \times N$ arrays using N^2 process
- Objectives: Learn parallel computing using CUDA & parallel decomposition of a problem
- Outcome: decomposed problem into sub-problems
learned how to use GPUs solved problems using threads on GPU cores.
- SW & HW: 64 bit CPU, 4GB RAM, CUDA toolkit, Nvidia, NVCC compiler Google colab.
- Theory: Dividing a computation into smaller computations & assigning them to different processors for parallel execution are two key steps in the design of parallel algorithm. The process of dividing a computation in smaller parts some or all which may potentially be executed in parallel is called decomposition.

Tasks are programmer defined units of computation into which the main computation is subdivided by means of decomposition.

Simultaneous execution of multiple tasks is the key to reducing time required to solve the entire problem.

Tasks can be of arbitrary size, but once defined they are regarded as indivisible units of computation.

The tasks into which a problem is decomposed may not all be the same size.

In addition of two vectors, we have to add i^{th} element from 1^{st} array with i^{th} element of 2^{nd} array to get i^{th} element of result.

This allocation can be allocated to a distinct thread. The same thing is done for vector product matrix.

Using CUDA, vectors can be added using

- n blocks, 1 thread/block
- 1 block, n threads
- m blocks, n threads/block
- Similarly the product of vector & matrix will result in a $1 \times n$ vector containing the

result of multiplication

- The product of two matrices will result in matrix of resultant.

CUDA Kernels & Threads

The fundamental part of a CUDA code is the kernel program.

Kernel is the function that can be executed in parallel in the GPU device.

A CUDA kernel is executed by an array of CUDA threads. All threads run the same code.

Each thread has an id that it uses to compute memory address & make control decisions. CUDA organizes thousands of threads into a hierarchy of a grid of thread blocks. A grid is a set of thread blocks that can be processed on a device in parallel.

A thread block is a set of concurrent threads that can cooperate among themselves through synchronization barriers & access to a shared memory space private to the block. Each thread is given an ID unique within the block each block has a unique ID within a grid.

~~kernel~~

- conclusion: successfully implemented & executed vector & matrix operations parallelly using CUDA.